

PRQL

a simple, powerful, pipelined SQL replacement

Aljaž Mur Eržen

Compiler developer @EdgeDB

August 30, 2023

```
from albums
filter album_id > 100
sort albums.title
take 10
join artists (==artist_id)
select {
    albums.album_id,
    albums.title,
    f"Artist name: {artist.name}",
}
```

Why?

There are transtion costs!

Overview

- 1 Flaws of SQL**
- 2 Language for relations**
- 3 Compiling queries**
- 4 PRQL, the project**

A deef dive into

Flaws of SQL

Origins of the relational model

1970, Edgar F. Codd: abstraction over data storage

→ Tuple relational calculus

1974, Donald D. Chamberlin & Raymond F. Boyce: SEQUEL

→ Not a “proper” programming language

Not really composable

```
SELECT album_id, COUNT(*) AS track_count  
FROM tracks GROUP BY album_id
```

Not really composable

```
SELECT i.album_id, i.track_count, a.artist_id
FROM (
    SELECT album_id, COUNT(*) AS track_count
    FROM tracks GROUP BY album_id
) AS i
JOIN albums a USING (album_id)
```


Patched syntax

```
SELECT  
    SUM(total)  
FROM  
    invoices
```

Patched syntax

```
SELECT
    total / SUM(total) OVER () AS normalized_total
FROM
    invoices
```

Patched syntax

```
SELECT DISTNICT name  
FROM invoices
```

Patched syntax

```
SELECT EVALUATE TYPE-EMPLOYEE  
  WHEN "F"  
    MOVE "FULL TIME" TO EMP-TYPE-PR  
  WHEN "P"  
    MOVE "PART TIME" TO EMP-TYPE-PR  
  WHEN "C"  
    MOVE "CONSULTANT" TO EMP-TYPE-PR  
  WHEN OTHER  
    MOVE "INVALID" TO EMP-TYPE-PR
```

Patched syntax

Too much syntax

... but also ...

Not enough syntax

Name resolution

```
SELECT title AS title_alias  
FROM albums
```

Name resolution

```
SELECT title AS title_alias  
FROM albums  
WHERE title_alias LIKE 'Do I Wanna %'  
GROUP BY title_alias  
ORDER BY title_alias
```

Name resolution

More rules:

- ORDER BY positionals
- Correlated subqueries
- LATERAL

Relations vs scalars

```
SELECT * FROM table
```

```
SELECT count(*) FROM table
```

Relations vs scalars

```
SELECT emp_id FROM emp WHERE role = 'manager'
```

Relations vs scalars

```
SELECT *  
FROM emp  
WHERE emp_id = (  
    SELECT emp_id FROM emp WHERE role = 'manager'  
)
```

Relations cannot be ordered

```
SELECT * FROM albums ORDER BY title
```

Relations cannot be ordered

```
SELECT
    *,
    ... AS my_col
FROM (
    SELECT * FROM albums ORDER BY title
) inner
```

Relations cannot be ordered

```
SELECT
    *,
    ROW_NUMBER()
        OVER (ORDER BY artist_id) AS my_col
FROM (
    SELECT * FROM albums ORDER BY title
) inner
```

Relations cannot be ordered

SELECT returns an **ordered set**

FROM pulls-in a **set**

Relations cannot be ordered

```
SELECT
    *,
    ... AS my_col
FROM (
    SELECT *
    FROM albums
) inner
ORDER BY title
```


Relations cannot be ordered

```
SELECT
    *,
    ... AS my_col
FROM (
    SELECT * FROM albums ORDER BY title LIMIT 10
) inner
ORDER BY title
```

Identity of aggregation

```
SELECT SUM(cost) FROM expenses WHERE FALSE
```

Two possible behaviors: NULL or 0

Both valid

Identity of aggregation

“Every marble in this bag is black”
... but the bag is empty.

Ancient greeks say FALSE

Modern logic says TRUE

SQL says NULL

Identity of aggregation

Homomorphism of addition

$$\text{SUM}([1]) + \text{SUM}([4, 5]) = \text{SUM}([1, 4, 5])$$

$$1 + 9 = 10$$

Identity of aggregation

$$\text{SUM}([1]) + \text{SUM}([]) = \text{SUM}([1])$$

$$1 + ? = 1$$

$$\rightarrow \text{SUM}([]) = 0$$

identity of addition

Identity of aggregation

```
COUNT( [] )      = 0
ARRAY_AGG( [] )  = []
SUM( [] )        = 0
ANY( [] )        = false
EVERY( [] )      = true
STRING_AGG( [] ) = ''
```

Dialects

Differences in:

- syntax (TOP vs LIMIT)
- available functions
- available data types

Dialects

A class of languages

There is a standard

Slight diviations

Dialects

Different:

- priorities
- backward compatibility guarantees
- implementation limitations

Dialects

No clear & robust specification

Compilers could:

- adapt query to target database
- produce error early

Design of a new

Language for relations

Tuple relational calculus

Relation \sim a set of tuples

```
SELECT track_id, name, title  
FROM tracks, albums  
WHERE tracks.album_id = albums.album_id
```

Tuple relational calculus

Relation \sim a set of tuples

$T := \text{tracks}$

$A := \text{albums}$

$\pi_{\text{track_id, name, title}}(T * A)$

Data model

Functions

Transforms

Grouping

Nulls

Modern micro-features

The task of a query language

Leaky abstractions

A better database interface

prqlc and its IRs,

general architecture

how to use it

how to contribute